CAT 2

Cloud Computing

Alejandro Pérez Bueno

Nov 26, 2023

Table of Contents

Exercise 1

Basic VM configuration

First and foremost, it is necessary to create the VM and access it with SSH:

```
ssh -i .ssh/id_rsa root@84.88.58.81 -p 55000
```

In order to configure the VM with minimal configuration, one must do the following (as *root*):

```
apt update && apt upgrade -y                                    ①
hostnamectl set-hostname tux                                    ②
echo "127.0.0.1 tux.ccuoc.org tux" >> /etc/hosts                ③
apt install nginx apache2-utils jmeter -y                       ④
```

① Update the system.
② Set hostname to `tux`.
③ Change domain to `ccuog.org`.
④ Install `nginx`, `apache2-utils` and `jmeter`.

Now everything has been set up (see Figure 1).

Set up web service

The website I created is a web version of the first practice I submitted (I write them in `quarto`, a markdown tool that does wonders for these kinds of assignments). It is fairly complex as it has custom images, code formatting, sidebars and other non-textual content.

First, copy the website and its files to the `/var/www/html` directory:

```
cp README.html /var/www/html/index.html                         ①
cp -r README_files/ /var/www/html/                              ②
cp -r img /var/www/html/                                        ③
```

① Main *HTML* file.
② Extra files from quarto to make the website pretty.
③ Images from screenshots I made in the previous CAA.

Then start the `nginx` service:

```
systemctl enable --now nginx
```

At this point our website is up and running! We can check that it is working with `ab` (ApacheBenchmark) and with `JMeter` (see Figure 2).

> **Note**
>
> Note how in the command from Figure 2a I did not specify the IP of VM, but the domain. This is merely to prove that it is possible to open the website this way too.

You can also simply open a browser on your local machine and enter the IP of the Remote (see Figure 3).

Exercise 2

Basic VM configuration

> **Warning**
>
> The VM from the previous exercise has been deleted, so nothing from the previous exercise will be reused in this one.

Three VMs have to be configured:

1. A machine called `frontend` with a private and a public NIC.
2. A machine called `worker1` with a private NIC.
3. A machine called `worker2` with a private NIC.

On all of them, we have to do the following:

- Update packages and install `apache2`.
- Set hostname and the hosts file appropriately following the same scheme as in the prior exercise.
- Specify the private IPs and domain names for the other VMs in all three machines, as follows:

```
192.168.0.7 frontend
192.168.0.8 worker1
192.168.0.6 worker2
```

You can see how I have set up these VMs in OpenNebula (see @on-vms-ex2).

Frontend setup

Edit the configuration file `/etc/apache2/sites-available/000-default.conf` adding the following lines inside the `VirtualHost` section:

```
<Proxy balancer://mycluster>
  BalancerMember http://worker1:80
  BalancerMember http://worker2:80
  ProxySet lbmethod=byrequests # or ProxySet lbmethod=bytraffic
</Proxy>
ProxyPass /balancer-manager !
ProxyPass / balancer://mycluster/
ProxyPassReverse / http://frontend.ccuoc.org/
ProxyPreserveHost On
<Location /balancer-manager>
  SetHandler balancer-manager
</Location>
```

This configuration sets a balancer for the two websites hosted on the worker VMs with a certain balancer algorithm (installed as a module using `a2enmod`). It also enables the use of the `balancer-manager` extension.

Workers setup

Copy the different websites and their necessary files to each of the worker VMs' `/var/www/html` directory. For this I picked the following:

- For `worker1` VM: Same site from the previous website.
- For `worker2` VM: Similar site from another project.

Testing the results

Once all 3 VMs are up and running, let's make sure that they all three have the `apache` service started and updated to the latest changes made in the configuration files. To do so, we will restart the server:

```
systemctl enable apache2
systemctl restart apache2
```

Then, using `ab` like in the previous exercise we can benchmark the two balancing algorithms we are to compare:

- `byrequests`: distributes traffic evenly to every worker. In this case it would be `50/50`.
- `bytraffic`: distributes the workload based on the load of each worker.

We will compare the two algorithms by running the following command:

```
ab -n 100 -c 10 <FrontendPublicIP>/
```

- `ab`: The same apache utility used in exercise 1.
- `-n 100`: number of petitions made to the server.
- `-c 10`: number of *concurrent* requests made to the server at a time.

As can be seen in the screenshots from Figure 5, `bytraffic` is more efficient at distributing higher workloads, since fewer requests are lost when compared to the `byrequests` algorithm.

You can also verify that the balancer located at `http://<FrontendPublicIP>/balancer-manager` is working (see Figure 6), and that both sites are served in the frontend (see Figure 7).

Exercise 3

> **Warning**
>
> For the third exercise, I reused some of the previous configurations. Here is what has changed in the setup:
> - Recreated `frontend`, without setting up Proxy (I kept running out of space when installing `docker`).
> - Reused `worker1`.
> - Deleted `worker2`.

Basic VM configuration

Let's begin by installing `docker` on both VMs. According to the docker documentation, we must enter the following commands:

```
# Add Docker's official GPG key:
sudo apt-get update
sudo apt-get install ca-certificates curl gnupg
sudo install -m 0755 -d /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/debian/gpg | \
  sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg
sudo chmod a+r /etc/apt/keyrings/docker.gpg

# Add the repository to Apt sources:
echo \
  "deb [arch=$(dpkg --print-architecture) \
    signed-by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/debian \
  $(. /etc/os-release && echo "$VERSION_CODENAME") stable" | \
  sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt-get update
```

And then install `docker` and other related utilities:

```
sudo apt-get install docker-ce docker-ce-cli containerd.io \
  docker-buildx-plugin docker-compose-plugin
```

To set up a `swarm` in `docker`, you can do the following on the `frontend` V:

```
docker swarm init --advertise-addr 192.168.0.25
```

This gives you a token that can be used on the `worker` VM as follows:

```
docker swarm join --token <token> 192.168.0.25:2377
```

Annex

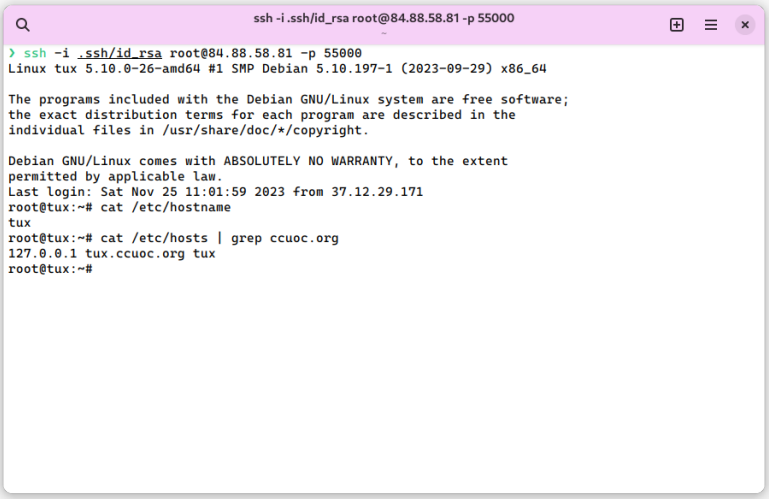

Figure 1: Minimal VM Configuration

‘‘

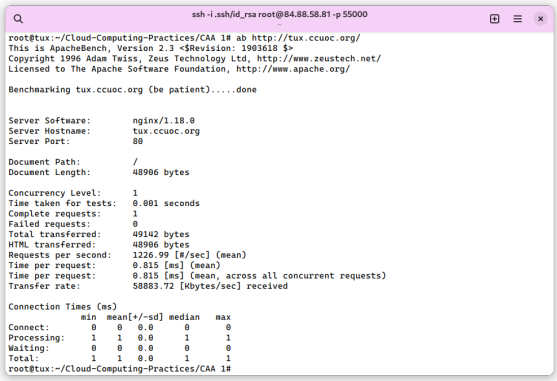

(a) `ab`

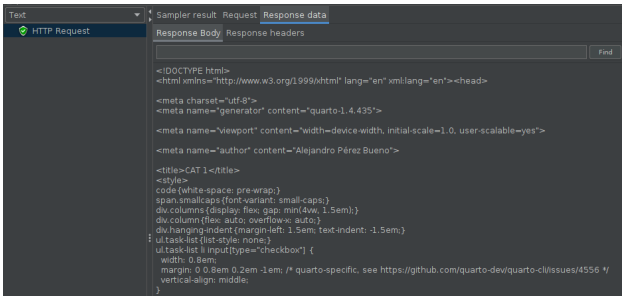

(b) `JMeter`

Figure 2: Test that the site is running with `ab` and `JMeter`

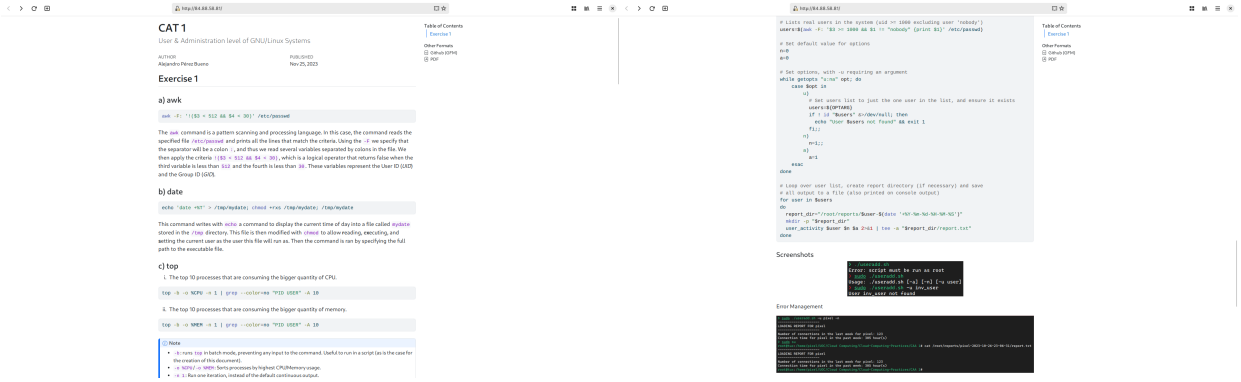

(a) Site top (b) Site image

Figure 3: Test that the site is running with a browser

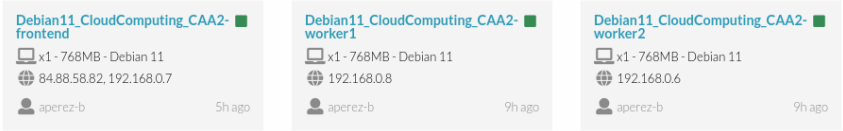

Figure 4: OpenNebula VMs

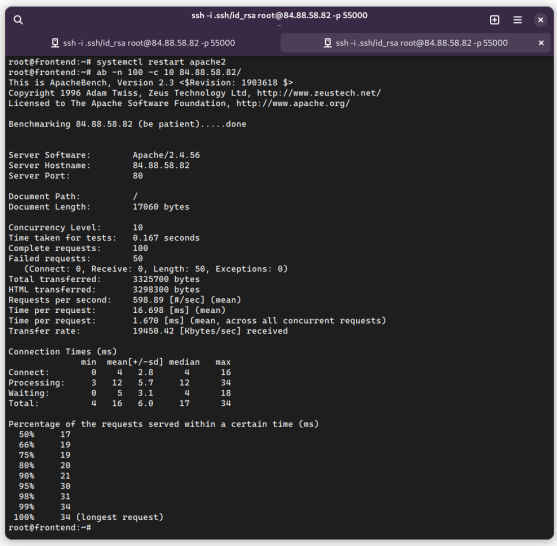

(a) `lbmethod=byrequests` - 50 failed requests

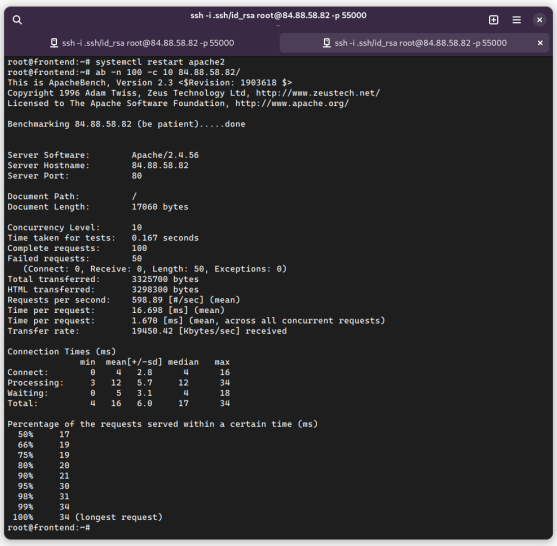

(b) `lbmethod=bytraffic` - 28 failed requests

Figure 5: Test the command `ab -n 100 -c 10 84.88.58.82` with two balancing algorithms

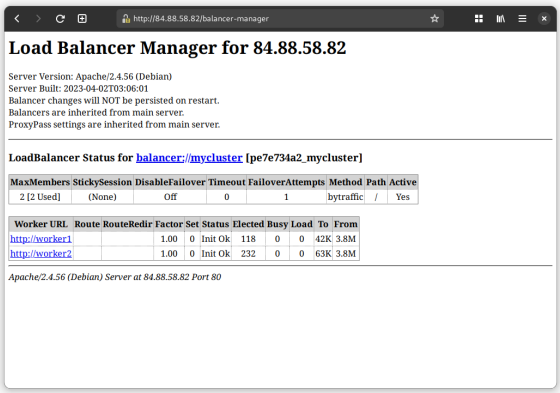

(a) Main `balancer-manager` page

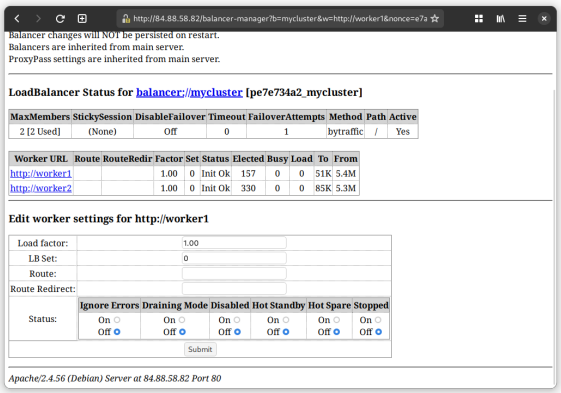

(b) Config for `worker1` in `balancer-manager`

Figure 6: `balance-manager` page

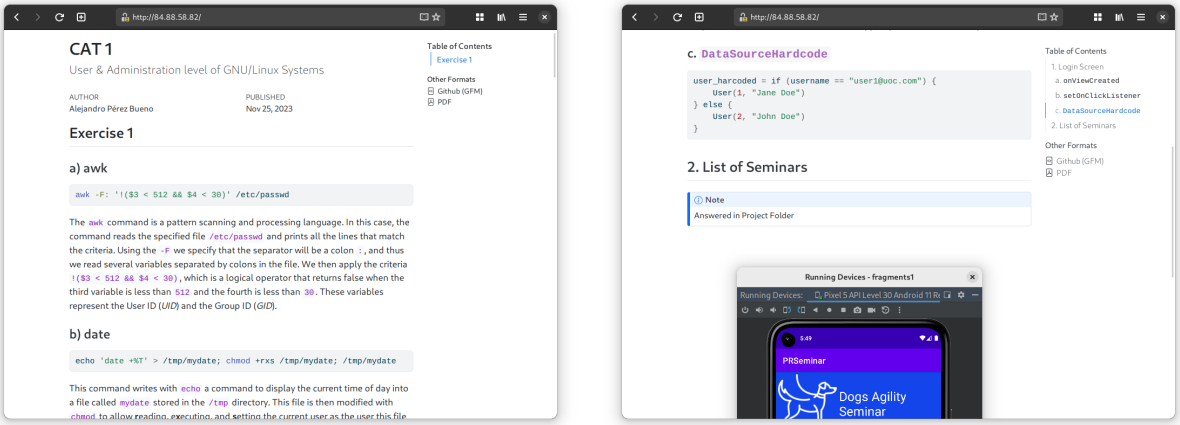

(a) Page from `worker1` accessed from `frontend` (b) Page from `worker2` accessed from `frontend`

Figure 7: Pages served by each worker balanced from the `frontend`